

▼ MODULE 4: Linear Regression - Assignment

In this assignment, we will show our understanding on how to apply the Linear Regression model to a selected dataset in Python.

Please, find the dataset on Canvas in the FILES section.

Description of the `advertising` data:

Advertising data sales (in thousands of units) for a particular product advertising budget (in thousands of dollars) for TV.

▼ STEP 1: Pre-process the dataset

In this step, you need to do the following:

- Import the necessary libraries
- Load the dataset
- Explore the dataset
- Summarize the dataset
- Analyze the dataset

▼ Import the necessary libraries

```
1 #Import the libraries  
2 import pandas as pd  
3 import numpy as np  
4 import seaborn as sn  
5 import matplotlib.pyplot as plt  
6
```

▼ Load the dataset

Comments:

```
1 #Load dataset  
2 df = pd.read_csv("/content/advertising.csv")
```

Explore the dataset

Note:

Here we can also check outliers, missing values, etc. You can apply all the DATA CLEANING skills that you learn in the Advanced Python course.

```
1 #Show the head of the data  
2 df.head()  
3
```

	TV	Sales	
0	230.1	22.1	grid
1	44.5	10.4	bar
2	17.2	12.0	
3	151.5	16.5	
4	180.8	17.9	

Next steps: [New interactive sheet](#)

```
1 #Show the shape of the data  
2 df.shape  
3
```

(200, 2)

```
1 #Get the info of the dataset with "advertising.info()"  
2 df.info()  
3
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 200 entries, 0 to 199  
Data columns (total 2 columns):  
 #   Column  Non-Null Count  Dtype     
---  --      --          --  
 0   TV      200 non-null    float64  
 1   Sales   200 non-null    float64  
dtypes: float64(2)  
memory usage: 3.3 KB
```

In the next step we will find the variable types which we can also find them using `.info()` as we did in the last step.

```
1 #Show the types of the variables
2 var_types = df.dtypes
3 var_types
4
```

```
0
TV    float64
Sales  float64
dtype: object
```

```
1 #Show the total null values in the data frame
2 df.isnull().sum()
3
```

```
0
TV    0
Sales  0
dtype: int64
```

▼ Summarize the data

```
1 #Use describe() for summarizing the data
2 #Round it to 2 decimal places
3 round(df.describe(),2)
4
```

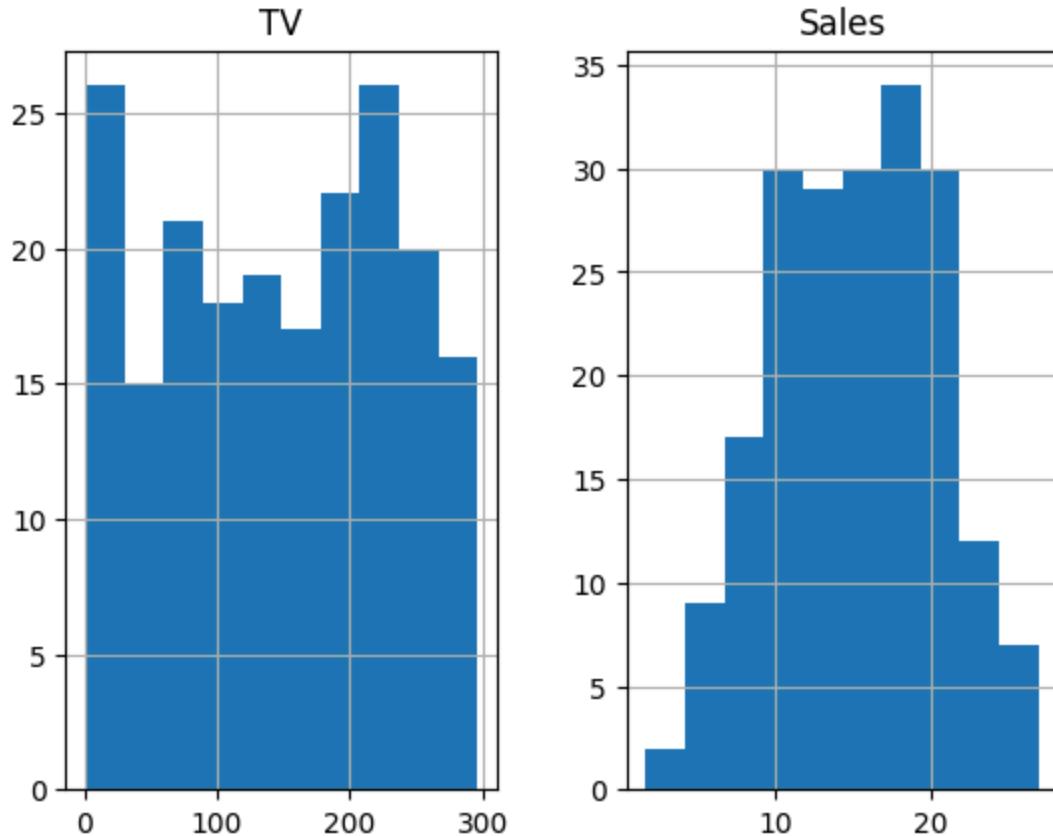
	TV	Sales	
count	200.00	200.00	
mean	147.04	15.13	
std	85.85	5.28	
min	0.70	1.60	
25%	74.38	11.00	
50%	149.75	16.00	
75%	218.82	19.05	
max	296.40	27.00	

▼ Analyze the data

Visualize the data

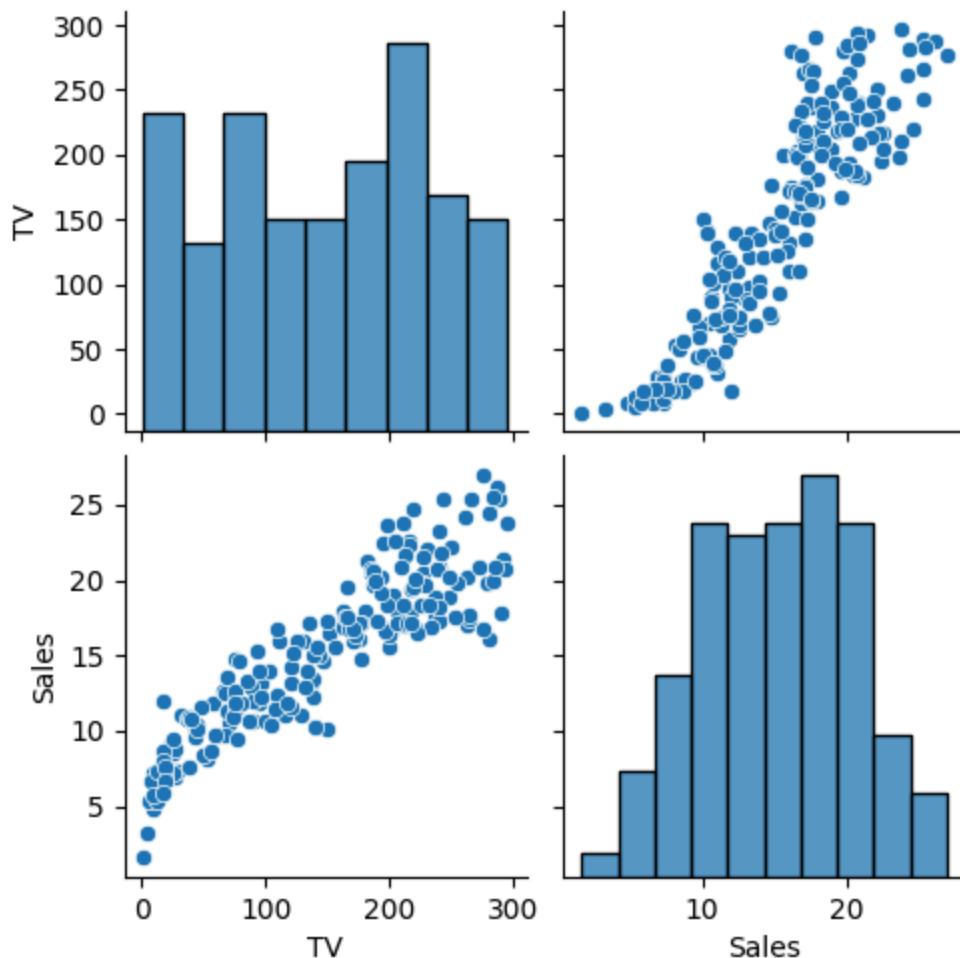
```
1 #Plot Histograms
2 hist_visual = df.hist(bins=10)
3 hist_visual
4
```

```
array( [[<Axes: title={'center': 'TV'}>,
         <Axes: title={'center': 'Sales'}>]], dtype=object)
```



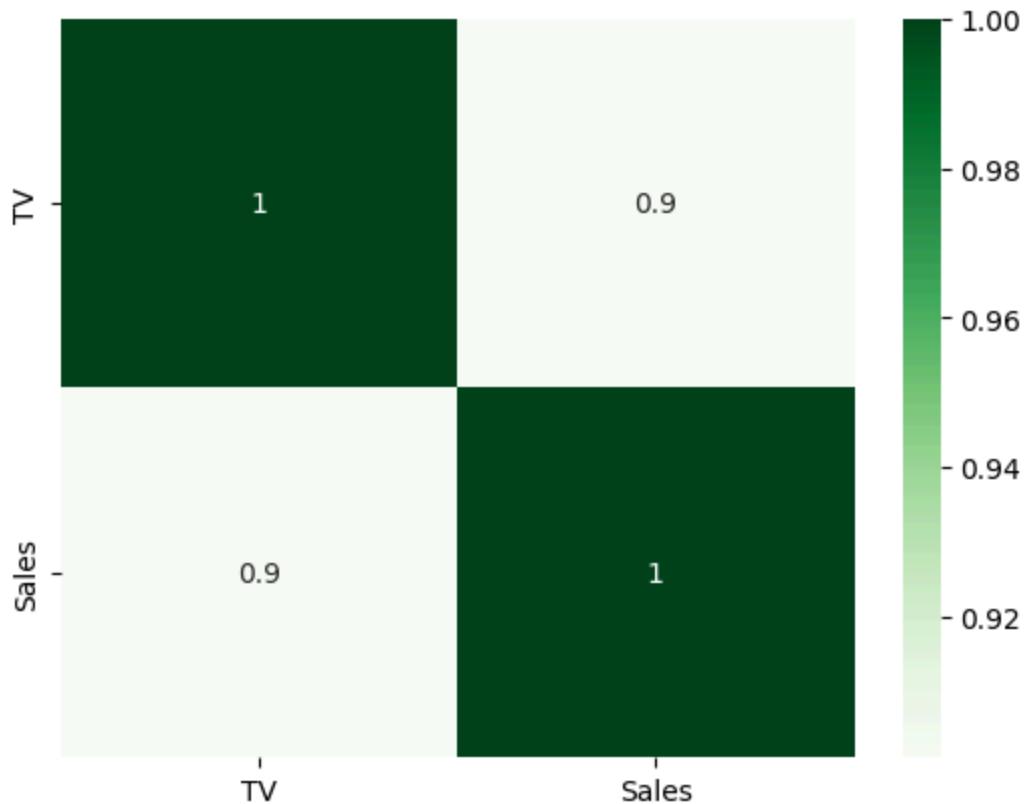
```
1 #Plot a pair plot to see the relationship between the variables
2 sn.pairplot(df)
3
```

<seaborn.axisgrid.PairGrid at 0x7bc7b9d62f90>



```
1 #Check correlation between TV and Sales
2 corr_matrix = df.corr()
3
4 #Heatmap
5 sn.heatmap(corr_matrix, annot=True, cmap= 'Greens')
6
```

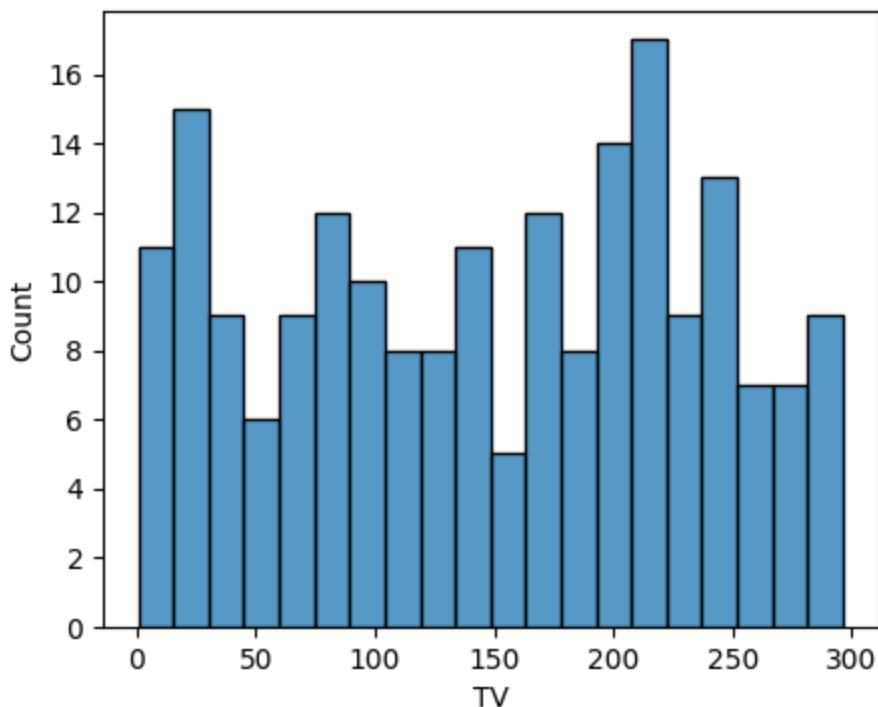
<Axes: >



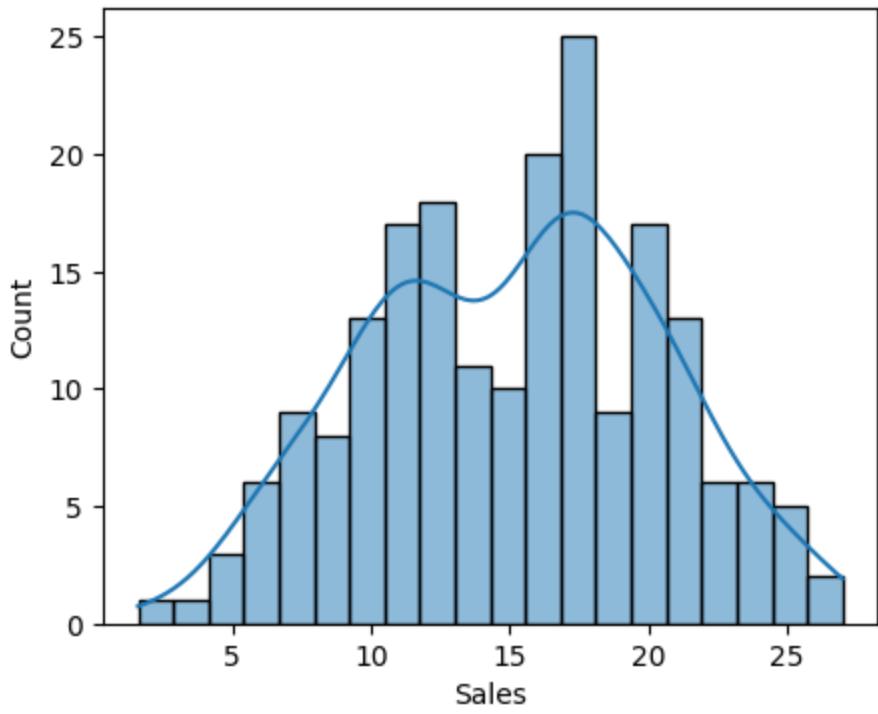
```
1 #Using the method "kendall" when using """.corr"""\n2 df.corr(method= 'kendall')\n3
```

	TV	Sales	grid icon
TV	1.000000	0.727994	bar chart icon
Sales	0.727994	1.000000	bar chart icon

```
1 #Do an outlier analysis for TV\n2 plt.figure(figsize=(5, 4))\n3 sn.histplot(df['TV'], bins=20)\n4 plt.show()\n5
```



```
1 #Do an outlier analysis for Sales(the target variable)
2 plt.figure(figsize=(5, 4))
3 sn.histplot(df['Sales'], bins=20, kde=True)
4 plt.show()
5
```

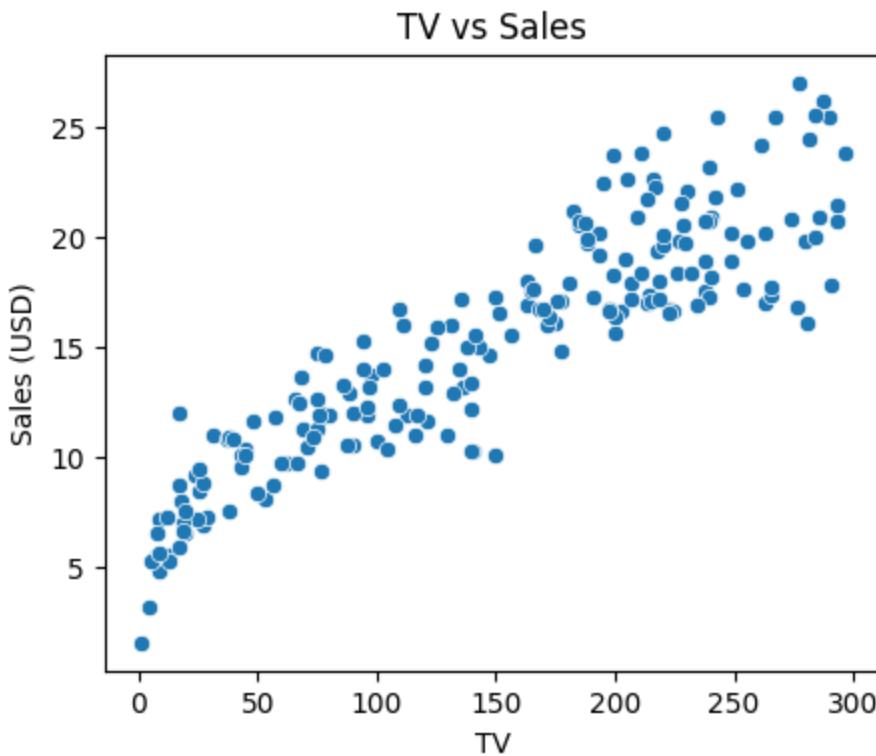


Question: Do you see any considerable outliers?

Answer: Yes there are a few

Plotting Salary vs. YearsExperience

```
1 #Plot a scatter plot to see how Sales is related to TV variable
2 plt.figure(figsize=(5,4))
3 sns.scatterplot(x='TV', y='Sales', data=df)
4 plt.title('TV vs Sales')
5 plt.xlabel('TV')
6 plt.ylabel('Sales (USD)')
7 plt.show()
```



Question: What can you see here about the correlation?

More TVs = More Sales

▼ STEP 2: Apply the Machine learning Model:

Here we will apply the ML model:

- Import the necessary libraries
- Build the model
- Display the results

Import the necessary libraries

```
1 #Import the library for splitting train/test data
2 from sklearn.model_selection import train_test_split
3
4 #Import for using the ML model
5 from sklearn.linear_model import LinearRegression
6 import statsmodels.api as sm #Linear regression with statsmodels
7
```

Buil the model

First we will prepare our dataframes for the x-array and y-array

```
1 #Create a X and Y data frames
2 X = df['TV']                      #used to make predictions
3 Y = df['Sales']                   #we'll predict this variable
```

Split the dataset into training and test dataset: Training 70% and test 30% sets

```
1 #Split the data into training and test sets
2 X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=
3
```

```
1 #Print the shape of the training and test data for X and Y
2 print("The shape of X_train is: ", X_train.shape)
3 print("The shape of X_test is: ", X_test.shape)
4 print("The shape of Y_train is: ", Y_train.shape)
5 print("The shape of Y_test is: ", Y_test.shape)
```

```
The shape of X_train is: (140,)
The shape of X_test is: (60,)
The shape of Y_train is: (140,)
The shape of Y_test is: (60,)
```

```
1 #The model requires a 2-D array
2 #Make the changes necessary for this
3 x_train = X_train.values.reshape(-1, 1)
4 x_test = X_test.values.reshape(-1, 1)
5
6 print("x_train shape is: ", x_train.shape)
7 print("x_test shape is: ", x_test.shape)
```

```
x_train shape is: (140, 1)
x_test shape is: (60, 1)
```

We build our linear regression model. We already import the necessary library.

```
1 #Set Linear Regression Model  
2 model = LinearRegression()  
3
```

Train the model with the training data. We use `.fit()` for that.

```
1 #Fit the training data  
2 response = model.fit(x_train, Y_train)  
3
```

Display Results

```
1 #Get the intercept and coefficient of the model  
2 intercept = response.intercept_  
3  
4 coeff = response.coef_  
5
```

```
1 #Print the intercept and coefficient of the model  
2 print("The intercept is: ", intercept)  
3 print("The coefficient is: ", coeff)
```

```
The intercept is: 6.928475121355413  
The coefficient is: [0.056007]
```

```
1 #Print the standard form using the intercept and coefficient  
2 print("Then, we have: y= %d + %d * x" %(intercept, coeff))  
3
```

```
Then, we have: y= 6 + 0 * x  
/tmp/ipython-input-3284058802.py:2: DeprecationWarning: Conversion of an  
print("Then, we have: y= %d + %d * x" %(intercept, coeff))
```